



# Java Servlets

## Basic Concepts and Programming

---

**Giuseppe Della Penna**

Università degli Studi di L'Aquila

*[giuseppe.dellapenna@univaq.it](mailto:giuseppe.dellapenna@univaq.it)*

*<http://people.disim.univaq.it/~dellapenna>*



# Notes to the English Version

*These slides contain an English translation of the didactic material used in the Web Engineering course at University of L'Aquila, Italy.*

*The slides were initially written in Italian, and the current translation is the first result of a long and complex adaptation work.*

*Therefore, the slides may still contain some errors, typos and poorly readable statements.*

*I'll do my best to refine the language, but it takes time.  
Suggestions are always appreciated!*

# Introduction to Servlets

- Servlets are special Java classes that run in specific web servers, called **servlet containers**.
- Servlets are exposed as a standard web resources (i.e., they can be referred to using a URL).
- Servlets act in the traditional *request/response* way, typical of the *server-side scripting*: when the user requests a servlet via the associated URL, the server activates it, executes it, and returns the result as the content of the resource.
- Java **Servlets APIs** allow one to program the servlet behavior without any knowledge of the server/client characteristics and transfer protocol.
- The code is standard Java code, which can make use of all the libraries and utilities for language, including the connection to all DBMS via JDBC, the use of XML through JAXP, etc..
- Servlets replace CGI providing a high degree of safety, versatility and abstraction to the programmer.

# Where and How to Run a Servlet

- To run a servlet, a particular server is required which can serve as **servlet container**, providing adequate support to their activation and execution.
- The most used *free* servlet container is **Tomcat** (Apache group).
- The traditional Apache server is not suitable for containing servlets, but it is possible to install it side-by-side to a Tomcat installation via a suitable *connector*.
- Tomcat is a *lightweight* container, and **implements only the basic technologies** for the development of web applications (*Servlets, JSP*). The complete **JEE Web profile** is available in more complex servers as *TomEE*, which is based on Tomcat, or *Glassfish*.

# Apache Tomcat Configuration

## Installation

- Apache Tomcat is available for all platforms (it is itself a Java program) and can be downloaded from <http://tomcat.apache.org/>.
- The installation on Windows and Unix is simplified by a fully automatic installation script.
  - On both platforms, you can choose to automatically start the server as a **service** (Windows) or **daemon** (UNIX), or **manually**.
  - The Java installation (or, better, the JRE) must be available to the installation script. To this aim, make sure the **JAVA\_HOME** environment variable is correctly set.
- Once executed, the default Tomcat instance responds on port **8080**.
- Via the url <http://localhost:8080/manager/> you can configure the server through a web application. and monitor the status of the web applications running in the server.
- To access such administrative applications, you should first **create a user with administrative privileges**, adding to the **conf/tomcat-users.xml** file a line like the following

```
<user username="admin" password="adminadmin" roles="admin,manager"/>
```

# Apache Tomcat Configuration

## Creating a new context

- Web applications are executed in **contexts**. In general, each context corresponds to a particular directory configured on the server and associated with a specific URL.
- To manually create a new context it is sufficient to create a subdirectory in the **webapps** directory of Tomcat. The context name will be the one of the directory.
- At this point, to test the new context, you can insert a plain html file in the directory and try loading the URL <http://localhost:8080/PATH/FILENAME>, where path is the name of the context. For example <http://localhost/project/index.html>
- In order to make a fully functional web application, we also need to prepare a special subdirectory structure in the context, and write some configuration files. The main elements of this structure are discussed below.
- However, the recommended installation method for web applications is to use an IDE (e.g., Netbeans) to create the application and package it in a war file (Web ARchive), which can then be copied directly into the webapps Tomcat directory.

# Apache Tomcat Configuration

## Structure of a context

- Context directories have a particular base structure which enables the server to access dynamic (servlets, JSP) and static (html, css, images, etc..) resources. In particular:
  - The WEB-INF subdirectory contains some configuration files, including the *web application deployment descriptor* (web.xml).
  - The WEB-INF/classes subdirectory contains the Java classes of the application, including servlets.
    - Following the Java conventions, individual classes should be placed in a directory structure corresponding to their package name.
  - The WEB-INF/lib subdirectory contains the JAR libraries used by the application, including the third party ones, such as JDBC drivers.
  - All other subdirectories of the context, including the root directory, will contain normal files as HTML pages, style sheets, images, or JSP pages.

# Apache Tomcat Configuration

## Adding a servlet to a context

- To make a Java class available as a servlet resource, you must configure its features through a file called **web application deployment descriptor**. This file, named **web.xml**, must be placed in the WEB-INF subdirectory of the context.
- A simple example of a descriptor is as follows:

```
<?xml version="1.0" encoding="ISO-8859-1"?>
<!DOCTYPE web-app PUBLIC "-//Sun Microsystems, Inc.//DTD
Web Application 2.3//EN" "http://java.sun.com/dtd/web-
app_2_3.dtd">
<web-app>
  <display-name>Progetto IW</display-name>
  <description>Progetto X</description>
  <servlet>
    <servlet-name>Servlet1</servlet-name>
    <description>Questa servlet implementa la funzione Y
  </description>
    <servlet-class>org.iw.project.class1</servlet-class>
  </servlet>
  <servlet-mapping>
    <servlet-name>Servlet1</servlet-name>
    <url-pattern>/funzione1</url-pattern>
  </servlet-mapping>
  <session-config>
    <session-timeout>30</session-timeout>
  </session-config>
</web-app>
```

- Each servlet is configured in a separate **<servlet>** element. The **<servlet-class>** element contains the full name of the class that implements the servlet.
- Each servlet must be mapped to a URL using a **<servlet-mapping>**. The specified **<url-pattern>** will compose the servlet URL as follows:

http://[server address]/[context]/[url  
pattern]



# The Servlet Base Classes

- The base for a servlet implementation is the **Servlet** interface, which is implemented by a set of base classes like **HttpServlet**. All the servlets will be derived (*extend*) from this abstract class.
- The other two base classes for the creation of a servlet are **ServletRequest** and **ServletResponse**.
- An instance of **ServletRequest** is passed from the context to the servlet when it is invoked, and contains all information related to the request: these include, for example, GET and POST parameters sent by the client, the server environment variables, *headers* and *payload* of the HTTP request.
- An instance of **ServletResponse** is passed to the servlet in order to return some content to the client. The methods of this class allow one to write to a *stream* which is then sent to the client, modify the HTTP response *headers*, etc..
- Other classes included in the servlet API, which we will not describe here, allow, e.g., to manage sessions.

# The Servlet Lifecycle

- The lifecycle of a servlet is marked by a sequence of calls made by container to particular methods of the Servlet interface.
  - **Initialization.** When the container loads the servlet, it calls its *init* method. Typically this method is used to establish connections to databases and prepare the context for subsequent requests. Depending on the context and/or container policies, the servlet can be loaded immediately when the server starts, or after the first request, etc..
  - **Service.** The client requests are handled by the container thorough calls to the *service* method. Concurrent requests correspond to executions of this method in separate threads. The implementation should therefore be thread safe. The service method receives user requests in the form of a *ServletRequest* and sends the response through a *ServletResponse*.
  - **Finalization.** When the container wants to remove/deactivate the servlet, it calls its *destroy* method. This method is usually used to close database connections, or dispose other persistent resources activated by the *init* method.
- The **HttpServlet** class specializes these methods for the HTTP communication. In particular, it contains two methods *doGet* and *doPost*, corresponding to the two most common HTTP verbs. The *service* method of *HttpServlet* class automatically redirects the client requests to the appropriate method.

# Writing a Servlet Class

- To write a simple servlet, you must create a class that extends **javax.servlet.http.HttpServlet**. The logic of the servlet is coded in the methods corresponding to the HTTP verbs.
  - **doGet** is called in response to GET and HEAD requests
  - **doPost** is called in response to POST requests
  - **doPut** is called in response to PUT requests
  - **doDelete** is called in response to DELETE requests
- All these methods have the same *signature*: they take a pair (*HttpServletRequest*, *HttpServletResponse*) and return *void*.
  - The **HttpServletRequest** and **HttpServletResponse** classes are specializations of *ServletRequest* and *ServletResponse* specific to the HTTP protocol.
- The **HttpServlet** class provides a default implementation of all these methods, which only generates an error 400 (BAD REQUEST)
- The servlet class contains other methods, such as *getServletContext*, which can be used to read a lot of information from the execution context of the servlet itself.
- To compile a servlet, the package **javax.servlet** must be included in the **CLASSPATH**. A copy of this library, called *servlet-api.jar* is present in the *common/lib* directory of Tomcat and in all the Java distributions.

# Writing a Servlet Class

## Example

```
package org.iw.project;

import javax.servlet.*;
import javax.servlet.http.*;

public class class1 extends HttpServlet {

    public void doGet(HttpServletRequest in,
        HttpServletResponse out) {
        //...
    }
}
```

- A servlet class extends the basic `javax.servlet.http.HttpServlet`
- For HTTP requests, the programmer should overwrite the appropriate methods: in this example, the `doGet` method is called to handle HTTP GET requests.

# Providing information on a Servlet

- It is possible, although not required, to provide information about servlets that can be used by the container.
- The information may include, for example the servlet author and the version number.
- For this purpose it is sufficient to override the *getServletInfo()* method on the Servlet interface, which returns *null* by default.
- The method takes no arguments and returns a string.

# Providing information on a Servlet

## Example

```
package org.iw.project;

import javax.servlet.*;
import javax.servlet.http.*;

public class class1 extends HttpServlet {

    public String getServletInfo() {
        return "Example servlet, version 1.0";
    }

    public void doGet(HttpServletRequest in,
        HttpServletResponse out) {
        //...
    }
}
```

- The string returned by **getServletInfo** the container will be used to provide information about the servlet.

# Initializing and Finalizing a Servlet

- The servlet initialization is accomplished in its **init** method, which has *ServletConfig* object as a parameter.
  - The first thing you should do is call the method **super.init()** passing it the *ServletConfig* parameter.
  - Then you can perform all the necessary initialization code, possibly by initializing the class fields with data that will be used by the service methods.
  - If the servlet has external initialization parameters (which are specified in a container-dependent way), you can read them through the *HttpServlet* method *getInitParameter*. This method takes as argument the name of the parameter and returns a string.
  - If the initialization has problems, you can throw a *ServletException* to report it to container.
- The servlet finalization is accomplished in its **destroy** method.
  - You must override this method only if there are things that you should do before the destruction of the servlet.

# Initializing and Finalizing a Servlet

## Example

```
package org.iw.project;

import javax.servlet.*;
import javax.servlet.http.*;

I
public class class1 extends HttpServlet {
    private int parameter1;

    public void init(ServletConfig c)
    throws ServletException {
        super.init(c);
        parameter1 = 1;
    }

    public String getServletInfo() {
        return "Example servlet, version 1.0";
    }

    public void doGet(HttpServletRequest in,
        HttpServletResponse out) {
        //...
    }
}
```

- The **init** method, after calling the same method of the upper class, proceeds with the servlet initialization.
- If there are initialization problems, it throws a **ServletException**.



# Reading the Request

- The **HttpServletRequest** object which is supplied to all the service methods contains information about the client's request.
  - The method *getParameter(String)* returns the value of a parameter included in the HTTP request. If the parameter can have more than one value, use *getParameterValues(String)*, which returns *the array* of all the values assigned to the given parameter.
    - These methods are not valid when you use the *multipart/form-data* encoding (file uploads)!
  - For GET requests, the method *getQueryString()* allows you to read the entire query string (not parsed).
  - For POST requests, methods *getReader()* and *getInputStream()* allow you to open a stream (text or binary) and directly read the request payload.
  - The method *getHeader()* allows one to read the contents of the HTTP request headers.
  - The method *getSession()* is used to manage sessions (see below).
  - Other methods allow to manage authentication, cookies, etc..
  - Methods inherited from class **ServletRequest**, finally, allow one to read information such request address (*getRemoteAddr*) or protocol (*getProtocol*).

# Reading the Request

## Example

```
package org.iw.project;

import javax.servlet.*;
import javax.servlet.http.*;
import java.io.*;

public class class1 extends HttpServlet {
    //...
    public void doGet(HttpServletRequest in,
        HttpServletResponse out) {
        String p1 = in.getParameter("p1");
    }
    public void doPost(HttpServletRequest in,
        HttpServletResponse out) {
        try {
            Reader r = in.getReader();
            //read the raw payload from r...
        } catch (Exception e) {
            e.printStackTrace();
        }
    }
}
```

- The doGet and doPost methods can read the parsed request parameters by **getParameter()** and **getParameterValues ()**.
- doGet can read the raw query string calling **getQueryString()**.
- doPost can read the payload of the message from the stream returned by **getReader() (text)** and **getInputStream() (binary)**.

# Writing the Response

- The **HttpServletResponse** object is provided to all the service methods and allows to build the reply to be sent to the client.
  - The method *setContentType(String)* is used to declare the return type (such as “text/xml”).
  - The method *SetHeader(String, String)* allows one to add further *headers* to the request.
  - The method *sendRedirect(String)* allows to redirect the browser to a new URL.
  - The methods *getWriter()* and *getOutputStream()* allow to open a channel, text or binary, to write the contents of the response. It should be called after setting the *content type* and writing any other *header*.
  - Other methods allow you to manage, for example, cookies.

# Writing the Response

## Example

```
package org.iw.project;

import javax.servlet.*;
import javax.servlet.http.*;
import java.io.*;

I
public class class1 extends HttpServlet {
    //...
    public void doGet(HttpServletRequest in,
        HttpServletResponse out) {
        out.setContentType("text/xml");
        try {
            Writer w = out.getWriter();
            w.write("pippo");
        } catch (Exception e) {
            e.printStackTrace();
        }
    }
}
```

- In the most common case, where you must return XHTML text to the client, it is sufficient to take the response **Writer** through the method *getWriter()* and use it to write all the text to be returned to the browser.
- Any other parameter of the response (in this case, the type) must be set *before opening* the output channel.

# The Sessions

- The concept of session is widely used in server-side programming. A session **associates state information to user requests**, allowing to circumvent the *stateless* feature of HTTP.
- Typically, sessions are associated to a **unique identifier** (*session identifier*) that is assigned to the user when he/she enters the web application (for example, after the login), then is sent together with each subsequent HTTP request, and finally it is invalidated when the user logs out or after a certain period of inactivity.
- The session identifier must be unique, and is usually generated with random algorithms based on the current date/time.
- To send the session identifier inside each request, we can typically use two approaches:
  - **Cookies:** Cookies are strings sent from the server to the browser. The browser automatically retransmits to the server the associated cookies together with each request, so there's no need for specific client-side code. The so-called *session cookies* are deleted from the browser when it closes, and contain the session identifier. Cookies are the easiest and widely adopted solution, but may be affected by the security policies of browsers (such as disabling cookies).
  - **URL rewriting:** session identifiers are written in the URL, as part of the path or, more commonly, as a GET parameter. This is the most generic and compatible solution, but requires the dynamic generation of all the website pages (each internal site link must be rewritten by introducing the current session identifier).

# Managing Sessions with Cookies

- In the servlet, creating and using a session through cookies is very simple. The session is managed by **HttpSession** objects. *Session variables* are simple strings that are associated with generic values of type **Object**.
- First, we must take a reference to the session object by requesting it through the **HttpServletRequest** method *getSession(boolean)*.
  - If the parameter to *getSession* is true, a new session is created when there is not a valid one already active. Otherwise, the function may return null.
  - Calling this method can change the response headers, so it must be called before starting the output.
- The **HttpSession** methods allow to manage the session:
  - *isNew()* returns true if the session has been just created: in this case, usually, its state variables must be initialized.
  - *getAttribute(String)* returns the object associated with the given name stored in the session.
  - *setAttribute(String, Object)* associates with the specified name the object passed as the second argument. In practice, it creates or updates the state variable given by the first argument using the value contained in the second argument.
  - *removeAttribute(String)* removes the given state variable.
  - *invalidate()* closes the session and deletes all the state information associated with it.

# Managing Sessions with Cookies

## Example

```
package org.iw.project;

import javax.servlet.*;
import javax.servlet.http.*;
import java.io.*;

I
public class class1 extends HttpServlet {
    //...
    public void doGet(HttpServletRequest in, HttpServletResponse out) {
        HttpSession s = in.getSession(true);
        if (s.isNew()) s.setAttribute("pagine", new Integer(1));
        int a = ((Integer)s.getAttribute("pagine")).intValue();
        s.setAttribute("pagine", new Integer(a+1));
        try {
            Writer w = out.getWriter();
            w.write("pagine visitate in questa sessione: "+a);
        } catch (Exception e) {
            e.printStackTrace();
        }
    }
}
```

- After each GET request to the servlet, if a session is not active, it is created and variable called "pages" initialized to 1 (note the use of the **Integer** class) is placed inside it.
- The number of pages visited during the session is then incremented and printed.

# Managing Sessions with Cookies

## Example

```
package org.iw.project;

import javax.servlet.*;
import javax.servlet.http.*;
import java.io.*;

public class class1 extends HttpServlet {
    //...
    public void doGet(HttpServletRequest in, HttpServletResponse out) {
        HttpSession s = in.getSession(true);
        s.setAttribute("user", in.getParameter("username"));
    }
}
```

- This simple servlet shows how to save in the session the value of the "username" parameter taken from the request (probably from a form).
- This is a typical operation performed at the end of a login process, to keep track of the user associated to the current session.



# Managing Sessions with URLs

- Servlets also have a semi-automated system to manage sessions via URL rewriting.
- In practice, in addition to the session management/creation/use code described in the previous slides, all you need is to transform any url that points to a resource in the same application using the method *encodeUrl(String)* of the object **HttpServletResponse**.
- The method `encodeURL` determines whether you must put the session identifier in the URL: if cookies are available, the URL is not altered.

# Managing Sessions with URLs

## Example

```
package org.iw.project;

import javax.servlet.*;
import javax.servlet.http.*;
import java.io.*;

I
public class class1 extends HttpServlet {
    //...
    public void doGet(HttpServletRequest in, HttpServletResponse
out) {
        HttpSession s = in.getSession(true);
        try {
            Writer w = out.getWriter();
            w.write(out.encodeUrl(in.getServletPath()));
        } catch (Exception e) {
            e.printStackTrace();
        }
    }
}
```

- In this example, a session is created (if necessary) and the current servlet's URL, rewritten by **encodeURL** to include the session identifier, is printed on the page.
- *If the browser supports cookies, the URL will not change.*

# Java and DBMS: the JDBC

- One of the most common operations in a web application is the **management of data stored in a database**.
- The data access in Java is done using the package **JDBC** (*Java DataBase Connectivity*). A typical use of the JDBC classes follows the following steps:
  - Make the **JDBC driver** for the DBMS in use available in the Java classpath.
  - Load the driver by making a reference to the class that implements it using *Class.forName* method
  - Proceed with the creation of a **Connention** object using the static method *getConnetion* of the **DriverManager** class.
    - The three parameters of the method are **the username and password** used to access the DBMS and the **JDBC connection string**, which specifies the address of the DBMS and the database to select: this string has a format that varies depending on the DBMS in use.
  - Create a **Statement** object on the connection, using the method *createStatement*.
  - Send an SQL query, as a string, to the DBMS through the created **Statement** and its method *executeQuery*.
    - The returned object, of **ResultSet** type, allows to browse the query results.
    - To send a query that returns no results, such as an insert or update statement, use the *executeUpdate* method. In this case, the return value is an integer representing the number of records processed by the query.
  - Once the results have been processed, free up the space reserved for them by calling the *close* method of **Statement**.
  - Finally, when the database is no more needed, close the corresponding connection by calling the *close* method of the **Connection**.
- All JDBC instructions, in case of error, raise exceptions derived from **SQLException**.

# Java and DBMS: the JDBC

## Example

```
import java.sql.*;

Class.forName ("com.mysql.cj.jdbc.Driver");

Connection con = DriverManager.getConnection(
    "jdbc:mysql://localhost/webdb?serverTimezone=Europe/Rome","
    user", "pass");

Statement stmt1 = con.createStatement();
ResultSet rs = stmt1.executeQuery(
    "SELECT * FROM test");
...
rs.close();
stmt1.close();

Statement stmt2 = con.createStatement();
int rc = stmt.executeUpdate("DELETE FROM test");
stmt2.close();

con.close();
```

- This example creates a connection to a *MySQL* database.
- The JDBC driver class, downloaded from the DBMS manufacturer's site, is *com.mysql.cj.jdbc.Driver*. Note: if you use a MySQL driver version prior to 8, the class name is *com.mysql.jdbc.Driver*.
- Warning: many servers have pre-installed drivers for common DBMS. However, they may not have the latest version, especially in the case of the MySQL driver version 8. In this case, add the driver as a library to your application!
- The connection string specifies the DBMS type (*mysql*) the DBMS listening point (*localhost*), and the database to select (*webdb*). It can also include other parameters specified as a query string. For the MySQL driver, starting from version 8, if a *timezone* is not set on the server, you will need to specify it using a parameter, as shown in the example.
- The connection also takes the username and password of the user to authenticate to the DBMS.
- First a selection query is executed via *executeQuery* and then a delete query through *executeUpdate*.

# Java and DBMS: the JDBC

## The ResultSet

- Through the **ResultSet** returned by the *executeQuery* method it is possible to read the columns of each record returned by a select query.
- The records must be read one at a time. At any time, the **ResultSet** points (via a *cursor*) to one of the records returned (*current record*).
- The values of the various fields of the current record can be read using the methods *GetX(column\_name)*, where X is the **Java base type** to extract (for example, *getString*, *getInt*, ...) and *column\_name* is the name of the field of the record to read.
- To move the cursor to the next record in the **RecordSet**, we use the *next* method. The method returns *false* when the records are ended.

# Java and DBMS: the JDBC

## ResultSet Example

```
import java.sql.*;

Class.forName ("com.mysql.cj.jdbc.Driver");

Connection con = DriverManager.getConnection(
    "jdbc:mysql://localhost/webdb?serverTimezone=Europe/Rome","user", "pass");

Statement stmt1 = con.createStatement();

ResultSet rs = stmt1.executeQuery(
    "SELECT * FROM test");

while (rs.next()) {
    System.out.println("nome = " + rs.getString("nome"));
}

rs.close();
stmt1.close();
con.close();
```

- In this example, we use a while loop to iterate through the results of a select query.
- For each record we print the value of the name field, of type string.

# Java and DBMS: the JDBC

## Limitations of the “standard” usage pattern

- In a *data-intensive* web application, where there are often many concurrent database accesses (many users may connect to the application simultaneously), the “standard” JDBC usage pattern presents considerable problems.
- In fact, **opening a database connection is usually a complex process**, since it requires
  - Loading drivers
  - Connecting to the Database Server
  - Authentication
- We need to limit the overhead due to these operations, to make web applications as fast as possible.

# Connection Pooling

- Connection pooling is a technique that allows to **simplify the procedures needed to open and close JDBC connections** using a **connections cache**, called the *connection pool*.
  - The pool maintains a set of connections to a database **already opened and initialized**.
  - When the application wants to connect, it takes a “ready” **connection from the pool**, and operates on it.
  - When the application closes the connection, it **actually remains open and is returned to the pool**, waiting to be reused for other requests.
  - The pool is initially filled with a certain number of connections. However, if the pool is empty (i.e., all its connections are in use) and new connections are requested, they are automatically created “on the fly”, enlarging the pool.
  - The connections left unused in the pool for too long can be closed automatically to free the corresponding DBMS resources.



# Connection Pooling

## Application server support

- Connection pooling support is built-in the most recent versions of JDBC, but it have to be implemented in third-party software (just like a JDBC driver).
- All the application servers provide an implementation of the connection pooling, and there are also external libraries used to add connection pooling support to any (non web) application.
- **Note: Each application server provides proprietary systems to configure connection pooling. We will see in particular how to use connection pooling with Tomcat.**

# Connection Pooling

## In Tomcat

- To configure connection pooling in Tomcat, we proceed as follows:
  - First, we configure the database connection in the web application context, modifying the server.xml (global server configuration) or, better, the context.xml (specific application configuration).
  - We add to the deployment descriptor (web.xml) a reference to the connection, which thus becomes an **application resource with type DataSource**.
  - In the code, we take a reference to the DataSource using the standard *Java Naming and Directory Interface* (JNDI).
  - We create a normal JDBC connection through the DataSource.
  - As usual, we close the connection when we don't need it anymore.

# Connection Pooling

## Application context configuration (context.xml) for Tomcat

```
<Context path="/Example_Database_Pooling">
  <Resource
    name="jdbc/webdb2"
    type="javax.sql.DataSource"
    auth="Container"
    driverClassName="com.mysql.jdbc.Driver"
    url="jdbc:mysql://localhost/webdb"
    username="website"
    password="webpass"
    maxActive="10"
    maxIdle="5"
    maxWait="10000"
  />
</Context>
```

- The connection is configured through a **Resource** element, which contains all its features, including the driver, username, password and connection string.
- The *type* must be specified as a **javax.sql.DataSource**. The *auth* attribute is set to "container".
- The *maxActive*, *maxIdle* and *maxWait* attributes are used to size the pool, indicating:
  - The highest number of connection that the pool will contain
  - How many unused connections are allowed in the pool
  - How long to wait for a new connection to become available
- **Warning:** The indicated JDBC driver must be copied in the lib directory of Tomcat. Simply copying it in the application WEB-INF/lib (as part of the deployment), will be it invisible to the class loader used for pooling!
- **Warning:** the class name and the structure of the connection string may change depending on the version of the driver used by Tomcat (see previous slides)

# Connection Pooling

## Deployment descriptor (web.xml) configuration

```
<web-app version="2.5"
  xmlns="http://java.sun.com/xml/ns/javaee"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://java.sun.com/xml/ns/javaee
    http://java.sun.com/xml/ns/javaee/web-app_2_5.xsd">
  ...
  <resource-ref>
    <res-ref-name>jdbc/webdb2</res-ref-name>
    <res-type>javax.sql.DataSource</res-type>
    <res-auth>Container</res-auth>
    <res-sharing-scope>Shareable</res-sharing-scope>
  </resource-ref>
  ...
</web-app>
```

- In the deployment descriptor we add a reference to the resource with a JNDI **resource-ref** element.
- The attributes *res-type* and *res-auth* reflect the *type* and *res* ones written in the **Resource** definition
- *res-sharing-scope* should generally be set to "Shareable"
- The *res-ref-name* attribute specifies the JNDI name used to access the resource in the code. The resources of type `DataSource` conventionally have a name that starts with "jdbc/".

# Connection Pooling

## Java code

```
try {
    //take a reference to the naming context
    InitialContext ctx = new InitialContext();
    //and lookup the Datasource in this context
    DataSource ds =
    (DataSource) ctx.lookup("java:comp/env/jdbc/webdb2");
    //connect to the database
    connection = ds.getConnection();
    //...use the connection...
} catch (NamingException ex) {
    //Exception raised if the requested resource does not exist
} catch (SQLException ex) {
    //JDBC standard exception
} finally {
    //at the end, the connection MUST be closed!
    try { connection.close(); } catch (SQLException ex) {}
}
```

- Create a *JNDI naming context* (**InitialContext**)
- Lookup the resource in the context. Note that the prefix “java:comp/env/” must be added to the resource name configured in the deployment descriptor. It can be useful to store such complete name in a web application initialization parameter.
- Cast the returned object to the actual type of resource (**DataSource**)
- Creates the **JDBC Connection** using the method **getConnection()** of the **DataSource**.
- After working on the connection, close it as usual, and it will be returned to the pool. **Warning: if the connection is not closed, it won't return in the pool!**

# Connection Pooling

## Resource Injection

```
class DatabaseService {  
    @Resource(name = "java:comp/env/jdbc/webdb2")  
    private DataSource ds;  
    //...  
    public void dbMethod() {  
        try {  
            connection = ds.getConnection();  
            //...use the connection...  
        } catch (SQLException ex) {  
            //JDBC standard exception  
        } finally {  
            //at the end, the connection MUST be closed!  
            try { connection.close(); } catch (SQLException ex) {}  
        }  
    }  
}
```

- **Resource injection** allows to skip the complex resources lookup process: indeed, Java itself will inject a reference to the *DataSource* in a user variable .
- The injection is usually performed on a class field, which **must have the correct type** (*DataSource*) .
- To perform injection, the field declaration must be preceded by the **@Resource** annotation, with the **name** parameter equal to the name of the resource to be injected

# ServletContextListener

- A *context listener*, can be useful to perform any “global application initialization” (i.e., not specific to a particular servlet).
- Objects that implement the interface **ServletContextListener** can be declared in the deployment descriptor (web.xml) as web application listeners.
- The two methods **contextInitialized** and **contextDestroyed** of these objects are called, respectively, when the web application starts and stops.
- These methods may perform any operation and change the **ServletContext** that will be then passed to all the servlets at runtime.

# ServletContextListener

## Example

```
public class ContextInitializer implements  
    ServletContextListener {  
  
    public void contextInitialized(ServletContextEvent sce) {  
        //initialize some global (context) variable...  
        sce.getServletContext().setAttribute("appID", 1);  
    }  
  
    public void contextDestroyed(ServletContextEvent sce) {  
  
    }  
}
```

```
<listener>  
  <listener-class>  
    it.univaq.f4i.iw.examples.ContextInitializer  
  </listener-class>  
</listener>
```

- This context listener initialize some context variables when the web application starts, and stores them as ServletContext attributes.
- Servlets can simply access the prepared DataSource with an instruction like **getServletContext ()**. **GetAttribute ("appID")**
- To activate the context listener, we simply add **the snippet below, which specifies the class name, to the web.xml.**



# Filter

- A *filter* can be used to modify «on the fly» the input (*HttpServletRequest*) and output (*HttpServletResponse*) data of a servlet.
- Objects that implement the **Filter** interface can be declared in the deployment descriptor (*web.xml*) as web application listeners.
- The two methods **init** and **destroy** of these objects are called, respectively, when the web application starts and stops.
- The **doFilter** method is invoked for each request to the servlets the filter is configured for, and receives the request and response objects and a *FilterChain*.
  - Filters can **change the request/response objects** with custom implementations to control the servlet input/output.
  - Filters must call the **doFilter** method of the *FilterChain*.

# Filter

## Example

```
public class EmailObfuscatorFilter implements Filter {  
    private FilterConfig config = null;  
    public void init(FilterConfig filterConfig) throws  
        ServletException {  
        this.config = filterConfig;  
    }  
    public void doFilter(ServletRequest request,  
        ServletResponse response, FilterChain chain) throws  
        IOException, ServletException {  
        chain.doFilter(request, response);  
    }  
    public void destroy() {  
        config = null;  
    }  
}
```

```
<filter>  
    <filter-name>emailfilter</filter-name>  
    <filter-class>EmailObfuscatorFilter</filter-class>  
</filter>  
<filter-mapping>  
    <filter-name>emailfilter</filter-name>  
    <url-pattern>*</url-pattern>  
    <dispatcher>REQUEST</dispatcher>  
</filter-mapping>
```

- This filter does nothing: it simply calls the *doFilter* method on the *FilterChain*.
- To activate the filter, we simply add the snippet below, which specifies the filter class name and the associated url patterns, to the web.xml.

# References

- **Servlet API**

<https://docs.oracle.com/javaee/7/api/javax/servlet/package-summary.html>

- **Servlet Tutorial**

<https://docs.oracle.com/javaee/7/tutorial/servlets.htm>

- **JDBC Tutorial**

<http://docs.oracle.com/javase/tutorial/jdbc>